

Final_Project_Template

April 6, 2024

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

The California Housing Data set from 1990 is a popular dataset often used for machine learning projects, particularly in regression analysis and geographic data visualization. It includes information about the housing in the California region, as recorded during the 1990 U.S. census.

Features:

- Median Income: The median income for households within a block of houses (measured in tens of thousands of USD).
- Housing Median Age: the median age of a house within a block; a lower number is a newer building.
- Total Rooms: The total number of rooms in the houses per block.
- Total Bedrooms: The total number of bedrooms in the houses per block.
- Population: The total number of people residing within a block.
- Households: The total number of households, a group of people residing within a home unit, for a block.
- Latitude: The block's latitude.
- Longitude: The block's longitude.
- Ocean Proximity: distance from the house to the ocean

Target Variable:

- Median House Value: This variable represents the median value of houses within a specific block, and it's what machine learning models aim to predict based on the other features in the dataset. Predicting the Median House Value helps in understanding housing market trends, assessing affordability, and planning economic and housing policies.

```
[2]: filename_data = 'data/housing.csv'
data = pd.read_csv(filename_data)
```

```
[3]: data.head(15)
```

```
[3]:    longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0      -122.23    37.88             41.0         880.0         129.0
```

1	-122.22	37.86	21.0	7099.0	1106.0
2	-122.24	37.85	52.0	1467.0	190.0
3	-122.25	37.85	52.0	1274.0	235.0
4	-122.25	37.85	52.0	1627.0	280.0
5	-122.25	37.85	52.0	919.0	213.0
6	-122.25	37.84	52.0	2535.0	489.0
7	-122.25	37.84	52.0	3104.0	687.0
8	-122.26	37.84	42.0	2555.0	665.0
9	-122.25	37.84	52.0	3549.0	707.0
10	-122.26	37.85	52.0	2202.0	434.0
11	-122.26	37.85	52.0	3503.0	752.0
12	-122.26	37.85	52.0	2491.0	474.0
13	-122.26	37.84	52.0	696.0	191.0
14	-122.26	37.85	52.0	2643.0	626.0

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY
5	413.0	193.0	4.0368	269700.0	NEAR BAY
6	1094.0	514.0	3.6591	299200.0	NEAR BAY
7	1157.0	647.0	3.1200	241400.0	NEAR BAY
8	1206.0	595.0	2.0804	226700.0	NEAR BAY
9	1551.0	714.0	3.6912	261100.0	NEAR BAY
10	910.0	402.0	3.2031	281500.0	NEAR BAY
11	1504.0	734.0	3.2705	241800.0	NEAR BAY
12	1098.0	468.0	3.0750	213500.0	NEAR BAY
13	345.0	174.0	2.6736	191300.0	NEAR BAY
14	1212.0	620.0	1.9167	159200.0	NEAR BAY

0.0.1 Getting to know the data

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
5   population              20640 non-null  float64
6   households              20640 non-null  float64
```

```

7  median_income      20640 non-null  float64
8  median_house_value  20640 non-null  float64
9  ocean_proximity    20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

```
[5]: data.describe()
```

```

[5]:      longitude      latitude  housing_median_age  total_rooms  \
count  20640.000000  20640.000000      20640.000000  20640.000000
mean    -119.569704    35.631861        28.639486    2635.763081
std       2.003532     2.135952        12.585558    2181.615252
min     -124.350000    32.540000         1.000000     2.000000
25%     -121.800000    33.930000        18.000000   1447.750000
50%     -118.490000    34.260000        29.000000   2127.000000
75%     -118.010000    37.710000        37.000000   3148.000000
max     -114.310000    41.950000        52.000000  39320.000000

      total_bedrooms  population  households  median_income  \
count  20433.000000  20640.000000  20640.000000  20640.000000
mean     537.870553   1425.476744    499.539680     3.870671
std     421.385070   1132.462122    382.329753     1.899822
min       1.000000     3.000000     1.000000     0.499900
25%     296.000000    787.000000    280.000000     2.563400
50%     435.000000   1166.000000    409.000000     3.534800
75%     647.000000   1725.000000    605.000000     4.743250
max    6445.000000  35682.000000   6082.000000    15.000100

      median_house_value
count      20640.000000
mean     206855.816909
std     115395.615874
min       14999.000000
25%     119600.000000
50%     179700.000000
75%     264725.000000
max      500001.000000

```

```
[6]: data.columns[data.isnull().any()]
data.isnull().sum()
```

```

[6]: longitude      0
latitude      0
housing_median_age  0
total_rooms    0
total_bedrooms  207
population     0
households     0

```

```

median_income      0
median_house_value  0
ocean_proximity    0
dtype: int64

```

Since there are null values in total bedrooms, we must perform operations on total bedrooms.

```
[7]: data[data["total_bedrooms"].isnull()]
```

```

[7]:      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
290      -122.16    37.77             47.0      1256.0             NaN
341      -122.17    37.75             38.0       992.0             NaN
538      -122.28    37.78             29.0     5154.0             NaN
563      -122.24    37.75             45.0       891.0             NaN
696      -122.10    37.69             41.0       746.0             NaN
...      ...      ...      ...      ...      ...
20267    -119.19    34.20             18.0     3620.0             NaN
20268    -119.18    34.19             19.0     2393.0             NaN
20372    -118.88    34.17             15.0     4260.0             NaN
20460    -118.75    34.29             17.0     5512.0             NaN
20484    -118.72    34.28             17.0     3051.0             NaN

```

```

      population  households  median_income  median_house_value  \
290         570.0       218.0         4.3750        161900.0
341         732.0       259.0         1.6196         85100.0
538        3741.0      1273.0         2.5762       173400.0
563         384.0       146.0         4.9489       247100.0
696         387.0       161.0         3.9063       178400.0
...      ...      ...      ...      ...
20267      3171.0       779.0         3.3409       220500.0
20268      1938.0       762.0         1.6953       167400.0
20372      1701.0       669.0         5.1033       410700.0
20460      2734.0       814.0         6.6073       258100.0
20484      1705.0       495.0         5.7376       218600.0

```

```

      ocean_proximity
290      NEAR BAY
341      NEAR BAY
538      NEAR BAY
563      NEAR BAY
696      NEAR BAY
...      ...
20267    NEAR OCEAN
20268    NEAR OCEAN
20372    <1H OCEAN
20460    <1H OCEAN
20484    <1H OCEAN

```

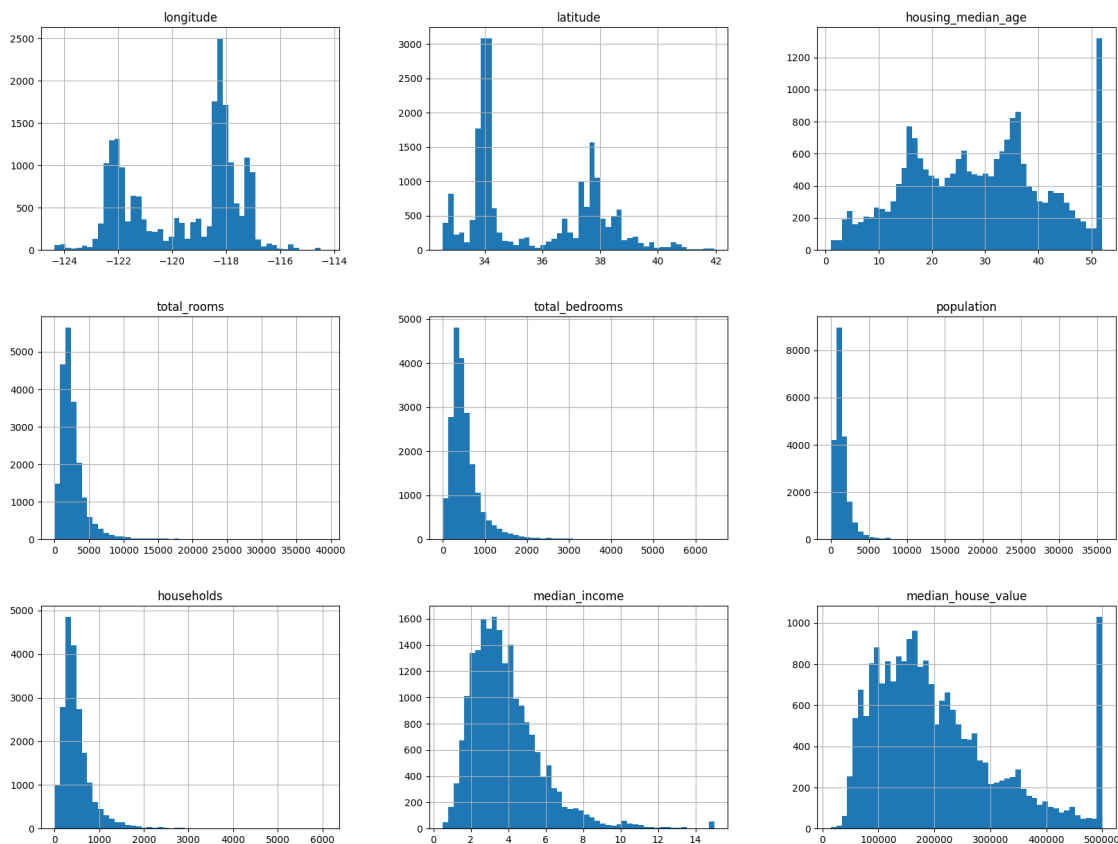
[207 rows x 10 columns]

```
[8]: data["total_bedrooms"] = data["total_bedrooms"].fillna(data["total_bedrooms"].
      ↪mean())
missing_total_bedrooms_values = data[data["total_bedrooms"].isnull()]
data[data["total_bedrooms"].isnull()]
```

[8]: Empty DataFrame
Columns: [longitude, latitude, housing_median_age, total_rooms, total_bedrooms, population, households, median_income, median_house_value, ocean_proximity]
Index: []

In total bedrooms, the average value of the existing values was assigned instead of the NaN values.

```
[9]: data.hist(bins=50, figsize=(20,15))
plt.show()
```



Since the ocean_proximity variable is an object, it was removed from the data because it caused problems in calculating the correlation matrix and did not have any effect on the target variable.

```
[10]: data = data.drop('ocean_proximity', axis=1)
```

One way of observing the strength of relationship between a feature and the target variable is to compute the correlation coefficient. The correlation coefficient ranges from -1 to 1. When it is close to 1, it means that there is a strong positive correlation. For example, the target tends to go up when the median_income goes up. When the coefficient is close to -1, it means that there is a strong negative correlation. Coefficients close to zero mean that there is no linear correlation.


```
[11]: corr_matrix=data.corr()  
      corr_matrix['median_house_value'].sort_values(ascending=False)
```

```
[11]: median_house_value    1.000000  
      median_income        0.688075  
      total_rooms          0.134153  
      housing_median_age    0.105623  
      households            0.065843  
      total_bedrooms        0.049454  
      population           -0.024650  
      longitude            -0.045967  
      latitude             -0.144160  
      Name: median_house_value, dtype: float64
```

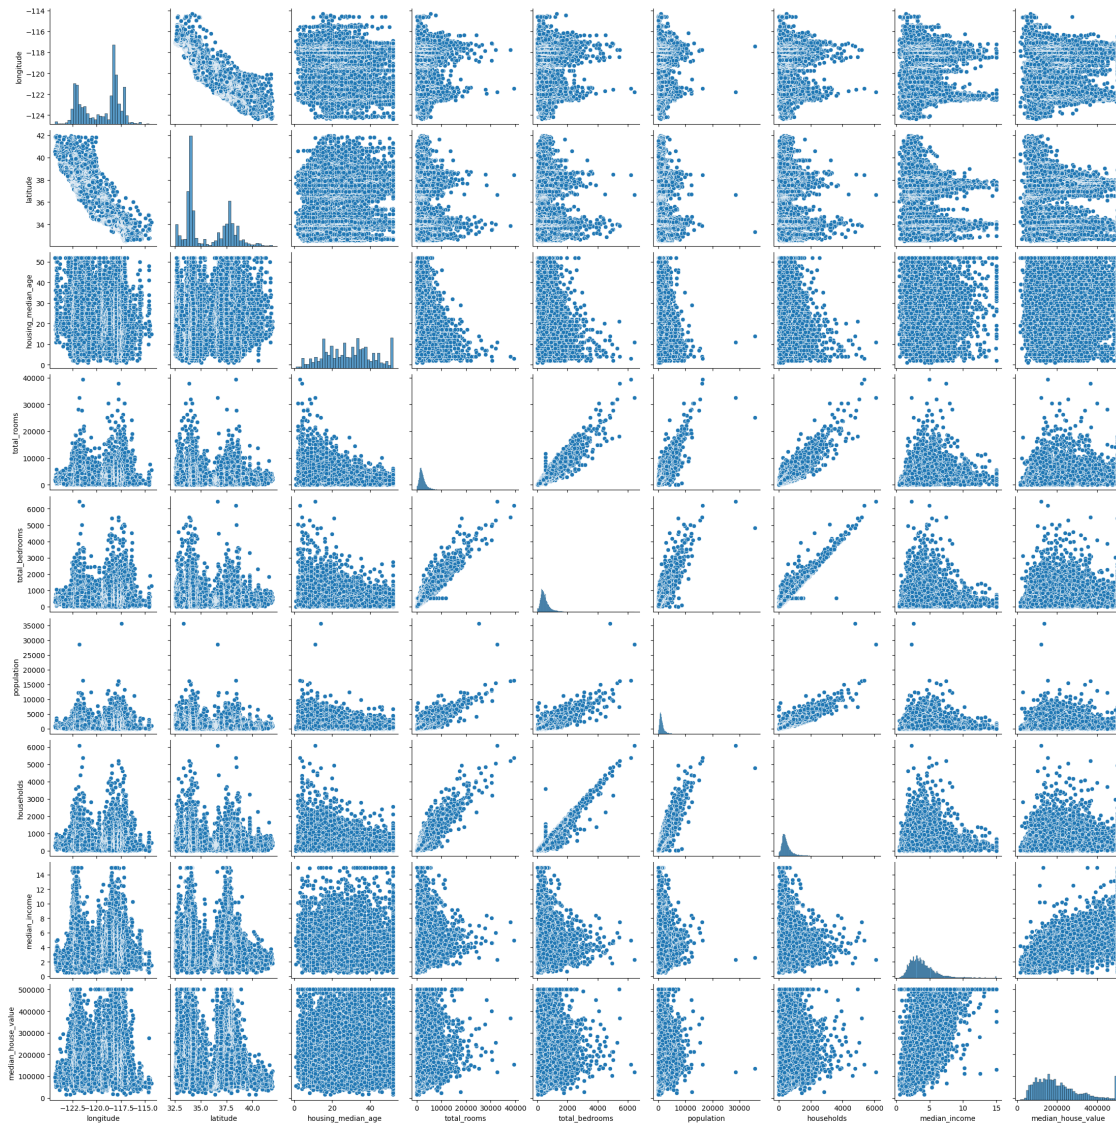
The relationship between the variables can also be observed visually, by plotting the scatter plot.

```
[12]: data.plot(kind='scatter', x='median_income', y='median_house_value', alpha=0.5)
```

```
[12]: <Axes: xlabel='median_income', ylabel='median_house_value'>
```



```
[13]: sns.pairplot(data)
plt.show()
```



All remaining features were observed to have a connection with the target variable.

0.0.2 Train/ Test Split

```
[14]: from sklearn.model_selection import train_test_split
train_set, test_set= train_test_split(data, test_size=0.2, random_state=42)
```

```
[15]: print("Number of examples in train set:", train_set.shape[0])
      print("Number of examples in test set:", test_set.shape[0])
```

Number of examples in train set: 16512

Number of examples in test set: 4128

```
[16]: trainX=train_set.drop('median_house_value', axis=1)
      trainY=train_set['median_house_value'].copy()

      testX=test_set.drop('median_house_value', axis=1)
      testY=test_set['median_house_value'].copy()
```

```
[17]: print(trainX.shape)
      print(trainY.shape)
      print(testX.shape)
      print(testY.shape)
```

(16512, 8)

(16512,)

(4128, 8)

(4128,)

0.0.3 Linear Regression

```
[18]: from sklearn.preprocessing import StandardScaler
      scaler = StandardScaler()
      trainX_scaled = scaler.fit_transform(trainX)
```

```
[19]: print(trainX_scaled)
```

```
[[ 1.27258656 -1.3728112  0.34849025 ...  0.76827628  0.32290591
   -0.326196  ]
 [ 0.70916212 -0.87669601  1.61811813 ... -0.09890135  0.6720272
   -0.03584338]
 [-0.44760309 -0.46014647 -1.95271028 ... -0.44981806 -0.43046109
    0.14470145]
 ...
 [ 0.59946887 -0.75500738  0.58654547 ...  0.28983345  0.07090859
   -0.49697313]
 [-1.18553953  0.90651045 -1.07984112 ...  0.30830275  0.15490769
    0.96545045]
 [-1.41489815  0.99543676  1.85617335 ...  1.04883375  1.94776365
   -0.68544764]]
```

```
[20]: from sklearn.linear_model import LinearRegression
      model=LinearRegression()

      model.fit(trainX_scaled, trainY)
```



```
[20]: LinearRegression()
```

```
[21]: from sklearn.metrics import mean_squared_error, r2_score
```

```
testX_scaled = scaler.transform(testX)
y_pred = model.predict(testX_scaled)
mse = mean_squared_error(testY, y_pred)
r2=r2_score(testY, y_pred)
print("Mean Squared Error on Test Set:", mse)
print("r2 score on Test Set:", r2)

y_pred = model.predict(trainX_scaled)
mse = mean_squared_error(trainY, y_pred)
r2=r2_score(trainY, y_pred)
print("Mean Squared Error on Train Set:", mse)
print("r2 score on Train Set:", r2)
```

```
Mean Squared Error on Test Set: 5052953703.90163
r2 score on Test Set: 0.6143987268246023
Mean Squared Error on Train Set: 4811134397.884197
r2 score on Train Set: 0.6400947924305294
```

0.0.4 Polynomial Regression

```
[22]: from sklearn.preprocessing import PolynomialFeatures
poly_features = PolynomialFeatures(degree=2, include_bias=False)
trainX_poly = poly_features.fit_transform(trainX_scaled)
print(trainX_poly.shape)
```

```
(16512, 44)
```

```
[23]: # Fit the polynomial regression model
model_poly = LinearRegression()
model_poly.fit(trainX_poly, trainY)
```

```
[23]: LinearRegression()
```

```
[24]: testX_poly = poly_features.transform(testX_scaled)
y_pred = model_poly.predict(testX_poly)
mse = mean_squared_error(testY, y_pred)
r2=r2_score(testY, y_pred)
print("Mean Squared Error on Test Set:", mse)
print("r2 score on Test Set:", r2)

y_pred = model_poly.predict(trainX_poly)
mse = mean_squared_error(trainY, y_pred)
```

```
r2=r2_score(trainY, y_pred)
print("Mean Squared Error on Train Set:", mse)
print("r2 score on Train Set:", r2)
```

```
Mean Squared Error on Test Set: 4577486232.73672
r2 score on Test Set: 0.6506826259019103
Mean Squared Error on Train Set: 3988584883.9903984
r2 score on Train Set: 0.7016270276689182
```

0.0.5 Ridge Regression

```
[25]: from sklearn.linear_model import Ridge
ridge_reg = Ridge(alpha=1.0)
ridge_reg.fit(trainX_scaled, trainY)
```

```
[25]: Ridge()
```

```
[26]: testX_scaled = scaler.transform(testX)
y_pred = ridge_reg.predict(testX_scaled)
mse = mean_squared_error(testY, y_pred)
r2=r2_score(testY, y_pred)
print("Mean Squared Error on Test Set:", mse)
print("r2 score on Test Set:", r2)

y_pred = ridge_reg.predict(trainX_scaled)
mse = mean_squared_error(trainY, y_pred)
r2=r2_score(trainY, y_pred)
print("Mean Squared Error on Train Set:", mse)
print("r2 score on Train Set:", r2)
```

```
Mean Squared Error on Test Set: 5052475906.315732
r2 score on Test Set: 0.6144351885395197
Mean Squared Error on Train Set: 4811135535.758371
r2 score on Train Set: 0.6400947073098889
```

```
[27]: plt.figure(figsize=(10, 6))
plt.scatter(trainY, y_pred, alpha=0.5)
plt.plot([min(trainY), max(trainY)], [min(trainY), max(trainY)], color='red')
plt.ylabel('Estimated Values')
plt.title('Real Values vs. Predicted Values')
plt.show()
```

